



Using the Couchbase C/C++ Client Library

Workshop Day 2

<https://github.com/dmaier-couchbase/cb-workshop-cpp>

<http://www.bit.ly/amadeus-cpp>

Before we begin



- Make sure that Couchbase Server is installed on the Dev Machine!



Document Modelling Basics

■ Java Script Object Notation

- Meta data
- Document Value

```
"meta" :  
{  
  "id" : "person::david",  
  "rev" : "1-0002bce00000000000",  
  "flags" : 0,  
  "expiration":0,  
  "type":"json"  
}  
  
"doc" :  
{  
  "type" : "person",  
  "uid":"david",  
  "firstname":"David",  
  "lastname":"Maier",  
  "birthday": 330004800000,  
  "email":"david.maier@couchbase.com"  
}
```

Normalization vs. De-Normalization

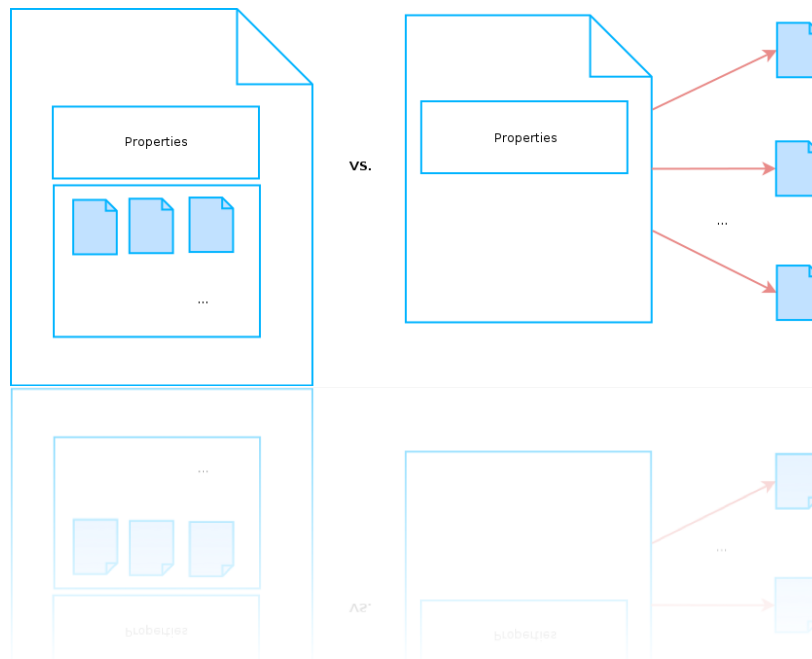


■ Normalized

- Uses key references for 1-many relationships
- Reduces data duplicates
- Smaller document size

■ De-Normalized

- Uses nested documents
- Aggregated view of data
- Allows atomic access
- No client side joins



Normalization vs. De-Normalization



DE-NORMALIZED

```
{
  "type" : "organization",
  "oid" : "CB",
  "name" : "Couchbase",
  "street" : "2440 West El Camino Real Suite 101",
  "city" : "Mountain View",
  "state" : "California"
  "employees" :
  [
    {
      "uid":"david",
      "firstname":"David",
      "lastname":"Maier",
      "birthday": 1402920000000,
      "email":"david.maier@couchbase.com"
    },
    ...
  ]
}
```

NORMALIZED

```
{
  "type" : "organization",
  "name" : "Couchbase",
  "street" : "2440 West El Camino Real Suite 101",
  "city" : "Mountain View",
  "state" : "California"
  "employees" : ["person::david", "person::perry", "person::dipti", ... ]
}
```

- Similar to sequences / auto-incrementing columns from the relational world
- Initialize and increment a counter value
- Use the counter as part of the key



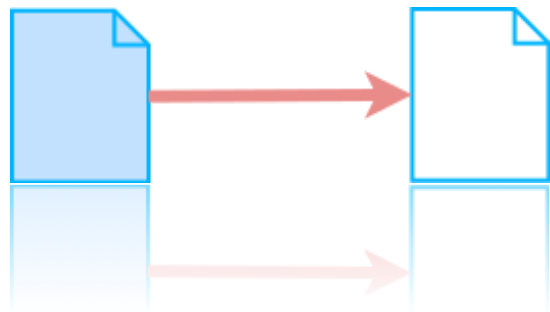
```
id = client.incr("count::person");
```

```
client.add("person::" + id, doc);
```

Reference Documents for Lookups



- Second document which references the primary one
- Needs to be maintained by the application



```
"email::david.maier@couchbase.com" : { "ref" : "person::david" }
```




Managing Connections

Exercise 7

Perform the following steps in order to install libcouchbase

- Perl needs to be installed

<http://developer.couchbase.com/documentation/server/current/sdk/c/start-using-sdk.html>

```
su root  
rpm -iv couchbase-release-1.0-2-x86_64.rpm  
yum install libcouchbase-devel libcouchbase2-bin
```



- The described setup procedure adds the the Couchbase package repository (/etc/yum.repos.d) and then installs the packages 'libcouchbase2-bin' and 'libcouchbase-devel'.

Get the Workshop Sources



Perform the following steps in order to check out the latest source code

- New installation

```
git clone https://github.com/dmaier-couchbase/cb-workshop-cpp.git
```

- Preinstalled workshop machine

```
cd ~/Git/cb-workshop-cpp  
git rebase
```

Before we begin



Open the documentation for libcouchbase!

- <http://developer.couchbase.com/documentation/server/4.5/sdk/c/st-art-using-sdk.html>
- Provided helper classes
 - CouchbaseDocument
 - CBCookie*
 - CBQStringConvert

Implement the following methods in CBDataSource:

- `void Connect(QString connectionString, QString password);`

Implement the following methods in CBDataSourceFactory:

- `static void Create(QString connectionString, QString password);`

Test your implementation by executing:

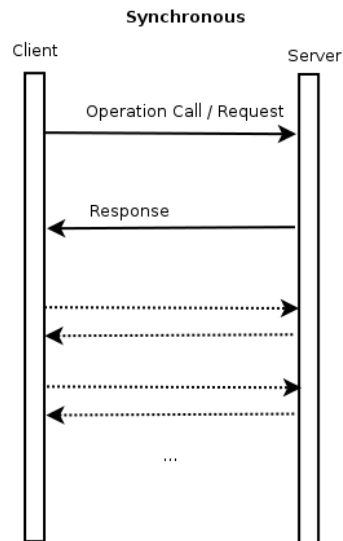
```
DemoCouchbaseConnect connectDemo;  
connectDemo.test();
```



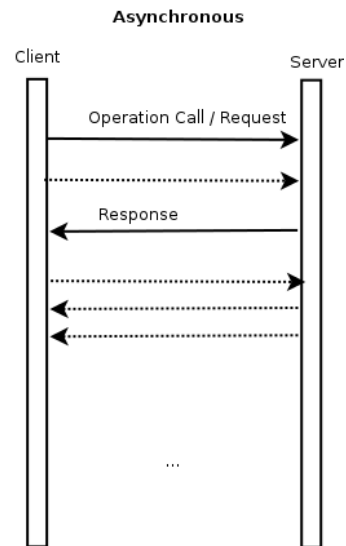
Understanding Non-Blocking I/O

in libcouchbase

- Libcouchbase is designed to use non-blocking I/O
 - Scheduled operations
- But `lcb_wait()` blocks by default
 - Waits for pending requests
 - Used for synchronous operation execution
- Callback functions are used
 - e.g. `storage_callback`



- External event loop integration
 - Provides mechanism to execute a callback function when a specific event occurs
 - e.g. libevent
- Asynchronous operation execution
- No need for `lcb_wait()`





Working with Documents

Exercise 8 - 11

Make sure that the travel-sample data is installed!
Implement the following methods in CBDataSource:

- CouchbaseDocument Get (QString key);

Test your implementation by executing:

```
DemoCouchbaseGet getDemo;  
getDemo.test();
```



Perform a Multi-Get



Make sure that the travel-sample data is installed!

Implement the following methods in CBDataSource:

- CouchbaseDocumentMap MultiGet(QStringList keys);

Test your implementation by executing:

```
DemoCouchbaseMultiGet multiGetDemo;  
multiGetDemo.test();
```

Implement the following methods in CBDataSource:

- bool Upsert(QString key, QString document)

Test your implementation by executing:

```
DemoCouchbaseUpsert upsertDemo;  
upsertDemo.test();
```

Delete a Document



Implement the following methods in `CBDataSource`:

- `bool Delete(QString key);`

Test your implementation by executing:

```
DemoCouchbaseDelete deleteDemo;  
deleteDemo.test();
```

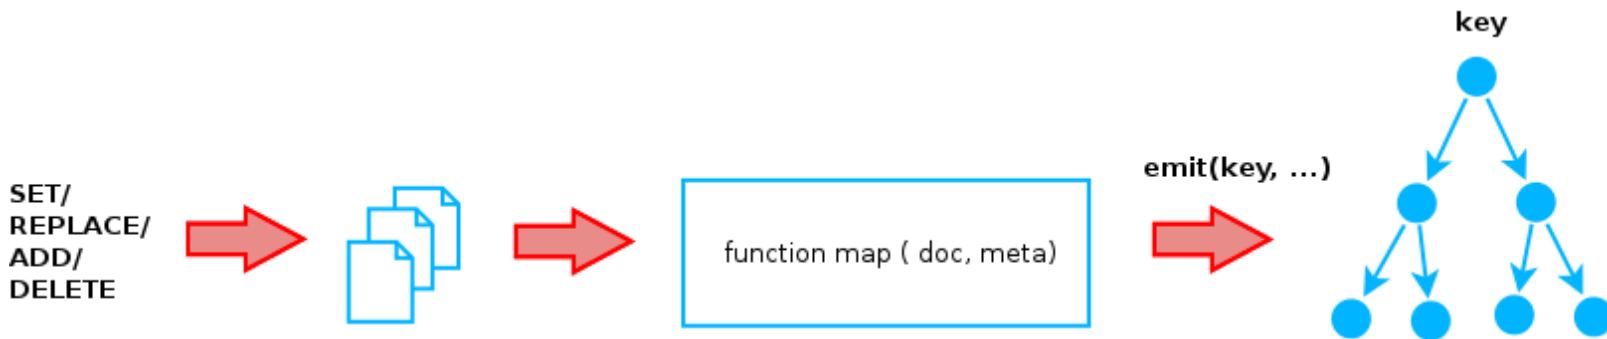


Querying via Views

Exercise 12

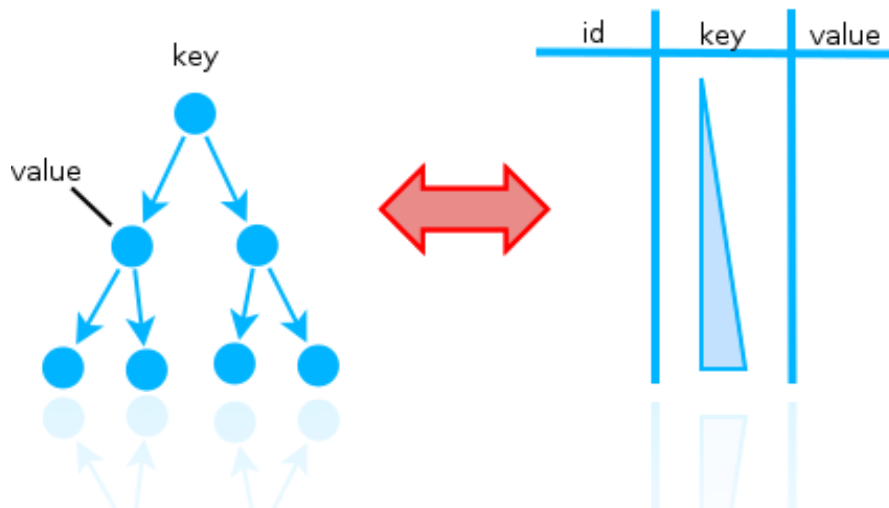
- Organized in Design Documents
- Incremental Map-Reduce
- Spread indexing load across nodes

Map	Reduce
Process, filter, map and emit a row	Aggregate mapped data Built in: _count, _sum, _stats



- Multiple roles

- A Primary Index to access all document id-s
- A Secondary Index as an alternative access path
- A View provides you an alternative view on your data



Create the View 'airports/by_name' !

Implement the following methods in CBDataSource:

- `CBQueryResult QueryView(QString designDocName,
 QString viewName, int limit=0, int skip=0);`

Test your implementation by executing:

```
DemoCouchbaseView viewDemo;  
viewDemo.test();
```





Querying via N1QL

Exercise 13

- Next generation, NoSQL query language
- SQL-like
 - WHERE
 - LIKE
 - GROUP
 - JOINS
- Powerful Extensions for JSON and hierarchical data structures
 - NEST
 - UNNEST
- Multiple access paths
 - Views
 - Global Secondary Indexes
 - Memory Optimized Indexes



N1QL Introduction - Joins



① Document Key: "customer802"

```
{
  "customer": {
    "ccInfo": {
      "cardExpiry": "2015-11-11",
      "cardNumber": "1212-1221-1121-1234",
      "cardType": "americanexpress"
    },
    "customerId": "customer802",
    "dateAdded": "2014-04-06T15:52:16Z",
    "dateLastActive": "2014-05-06T15:52:16Z",
    "emailAddress": "r_blonde@gmail.com",
    "firstName": "Richard",
    "lastName": "Blond",
    ...
    "postalCode": "05905",
    "state": "VT",
    "type": "customer"
  }
}
```

② Document Key: "purchase650"

```
{
  "purchases": {
    "customerId": "customer802",
    "lineItems": [
      { "count": 3,
        "product": "product55" },
      { "count": 4,
        "product": "product69" },
      ...
    ],
    "purchaseId": "purchase7049",
    "type": "purchase"
  }
}
```

③ Document Key: "purchase914"

```
{
  "purchases": {
    "customerId": "customer802",
    "lineItems": [
      { "count": 5,
        "product": "product55" },
      { "count": 3,
        "product": "product549" },
      ...
    ],
    "purchaseId": "purchase3648",
    "purchasedAt": "2013-11-07T15:52:38Z",
    "type": "purchase"
  }
}
```

```
1 SELECT c.emailAddress, count(p)
2 FROM purchases p
3 JOIN customers c
4 ON KEYS (p.customerId)
5 GROUP BY c.emailAddress;
```

N1QL Query Examples



```
SELECT airportname FROM `travel-sample` WHERE  
faa='LAX'
```

```
SELECT airportname FROM `travel-sample` WHERE  
faa='LHR'
```

```
SELECT faa as fromAirport,geo FROM `travel-sample`  
WHERE airportname = 'Los Angeles Intl' UNION  
SELECT faa as toAirport,geo FROM `travel-sample`  
WHERE airportname = 'Heathrow'
```

```
SELECT r.id, a.name, s.flight, s.utc, r.sourceairport,  
r.destinationairport, r.equipment FROM `travel-sample` r  
UNNEST r.schedule s JOIN `travel-sample` a ON KEYS  
r.airlineid WHERE r.sourceairport='LHR' AND  
r.destinationairport='LAX' AND s.day=6 ORDER BY a.name
```

```
SELECT airportname FROM `travel-sample` WHERE  
airportname LIKE 'Los An%'
```

Make sure that at least a Primary Index is created!

Also Double check that the Secondary Index on 'faa' is there!

Implement the following methods in CBDataSource:

- CBN1qlResult QueryN1ql(QString query);

Test your implementation by executing:

```
DemoCouchbaseN1ql n1qlDemo;  
n1qlDemo.test();
```





Error Handling and Logging

- Operations return `lcb_error_t` status code
- Check for
 - `err == LCB_SUCCESS`
- Error Codes
 - `<libcouchbase/error.h>`
- Examples
 - `LCB_KEY_EEXISTS`: Key already exists
 - `LCB_KEY_ENOENT`: Key does not already exist if replacing it
 - `LCB_ETIMEDOUT`: Transient error which indicates that something took too long
 - `LCB_ETMPFAIL`: Transient error which indicates that the server was too busy
 - `LCB_AUTH_ERROR`: Authentication error
 - `LCB_BUCKET_ENOENT`: Bucket does not exist

- LCB_LOGLEVEL environment variable
 - 1 – basic
 - 5 – verbose
- Programmatically
 - LCB_CNTL_CONLOGGER_LEVEL setting
 - console_log_level option in the connection string
- Log entry format

```
1ms [lo] {14780} [DEBUG] (lcbio_mgr - L:383) <localhost:11210> (HE=0xe56760)  
Creating new connection because none are available in the pool
```



A Sample Application

Exercise 14

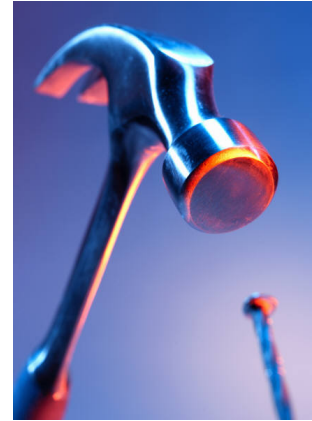
<https://github.com/dmaier-couchbase/cb-workshop-cpp/tree/master/TravelAppSample>

A Sample Application

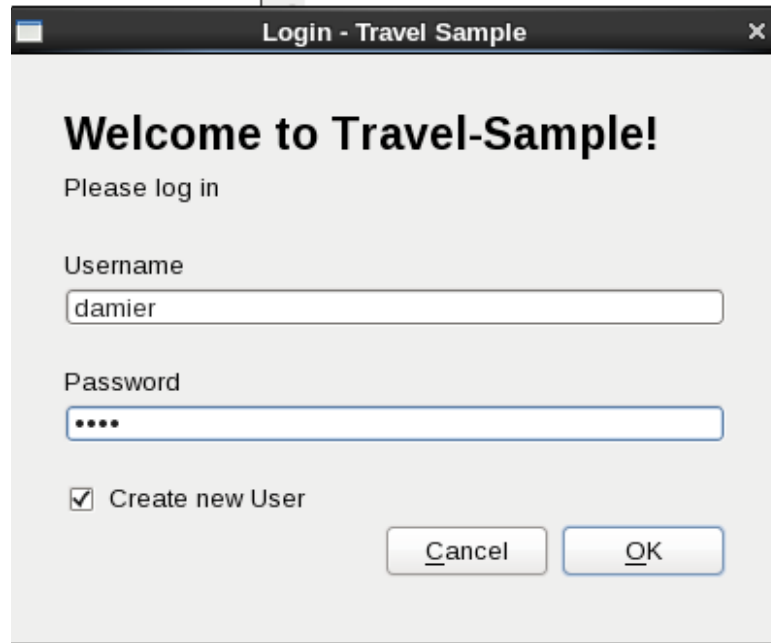


Inspect the full source code of the Travel-Sample application!
Run the Qt application!

- Search for a flight from 'LAX' to 'LHR'



A Sample Application



The screenshot shows a standard login window with a title bar, a welcome message, a login prompt, input fields for username and password, a checkbox for creating a new user, and 'Cancel' and 'OK' buttons.

Login - Travel Sample

Welcome to Travel-Sample!

Please log in

Username

damier

Password

....

☒ Create new User

Cancel OK

A Sample Application



Travel Sample

Couchbase Travel Sample C++/Qt

Logged in user: david

Flight Selection Available Flights Shopping Cart (0) Bookings (0)

Airport or City

From To

Los Angeles Intl

Heathrow

Travel Dates

Leave Return

Options

☒ Round Trip Travelers

Find Flights

A Sample Application



Travel Sample

Couchbase Travel Sample C++/Qt Logged in user: david

Flight Selection Available Flights Shopping Cart (0) Bookings (0)

Outbound Leg

Airline	Flight	Departure	From	To	Aircraft	Price
American Airl...	AA632	21:35:00	LAX	LHR	77W	\$850
American Airl...	AA951	23:12:00	LAX	LHR	77W	\$762
American Airl...	AA938	11:57:00	LAX	LHR	77W	\$780
American Airl...	AA275	03:18:00	LAX	LHR	77W	\$797
British Airways	BA353	03:19:00	LAX	LHR	744 388	\$718
British Airways	BA245	05:21:00	LAX	LHR	744 388	\$815

Add to cart

Inbound Leg

Airline	Flight	Departure	From	To	Aircraft	Price
American Airl...	AA208	04:27:00	LHR	LAX	77W	\$736
American Airl...	AA005	07:18:00	LHR	LAX	77W	\$806
American Airl...	AA323	14:41:00	LHR	LAX	77W	\$788
British Airways	BA330	13:19:00	LHR	LAX	744 388	\$736
British Airways	BA069	01:36:00	LHR	LAX	744 388	\$797
British Airways	BA433	18:36:00	LHR	LAX	744 388	\$797

Add to cart

A Sample Application



Travel Sample

Couchbase Travel Sample C++/Qt

Logged in user: david

Flight Selection Available Flights Shopping Cart (2) Bookings (0)

Airline	Flight	Departure	From	To	Aircraft	Price
American Airl...	AA275	03:18:00	LAX	LHR	77W	\$797
American Airl...	AA005	07:18:00	LHR	LAX	77W	\$806

remove selected item

Total Price: **\$1603**

Book now



Q&A

<http://docs.couchbase.com/developer/c-2.4/c-intro.html>



Specific Use Case Presentation

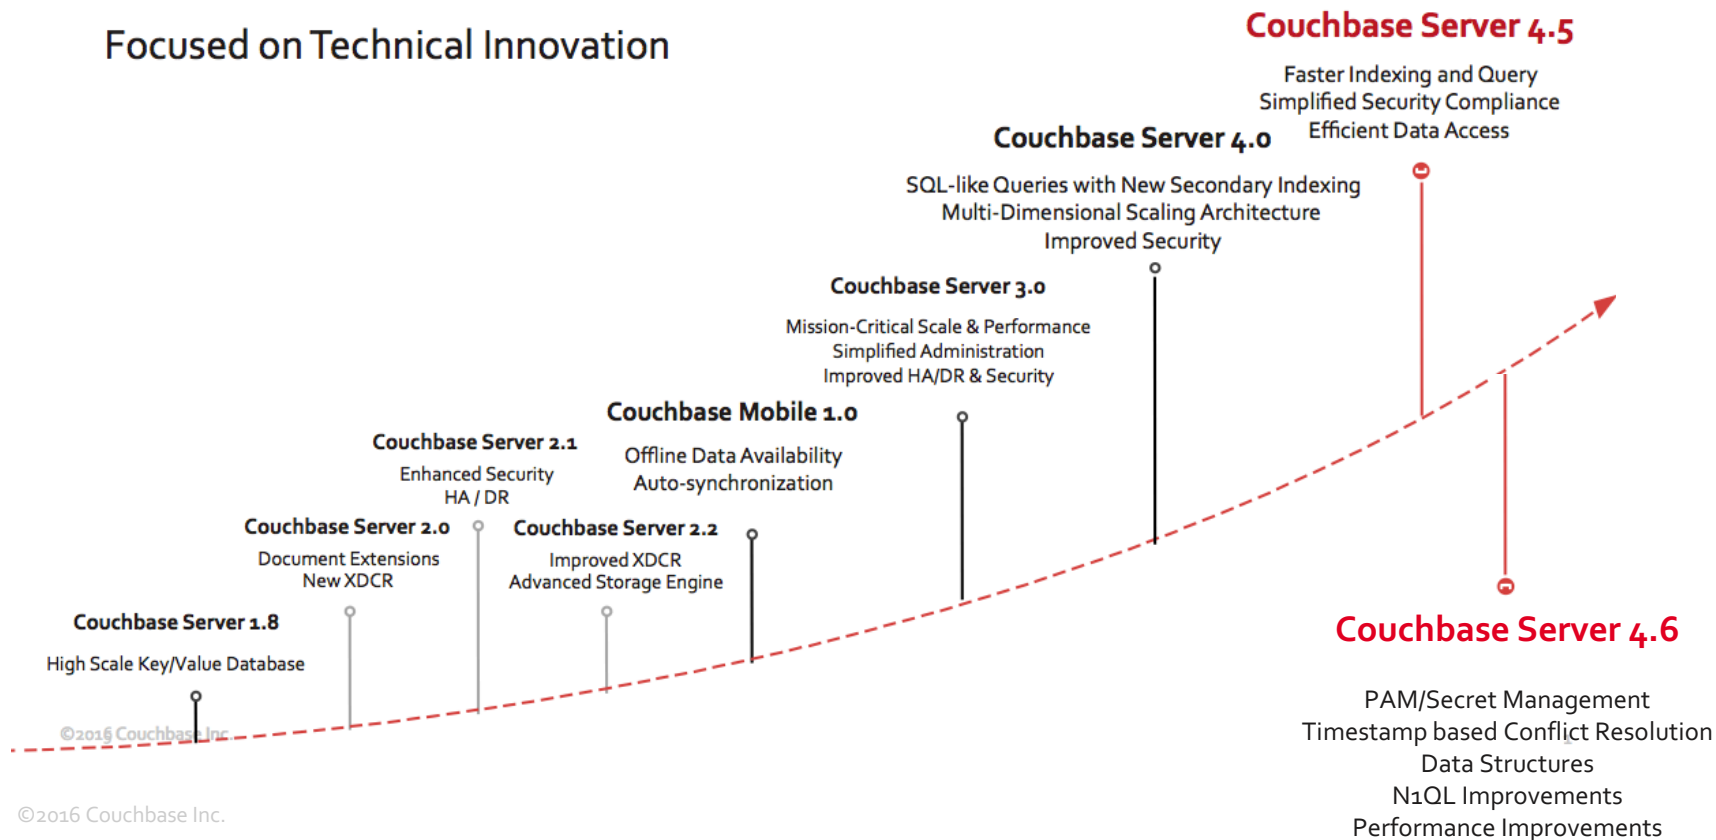


What's new in 4.x?

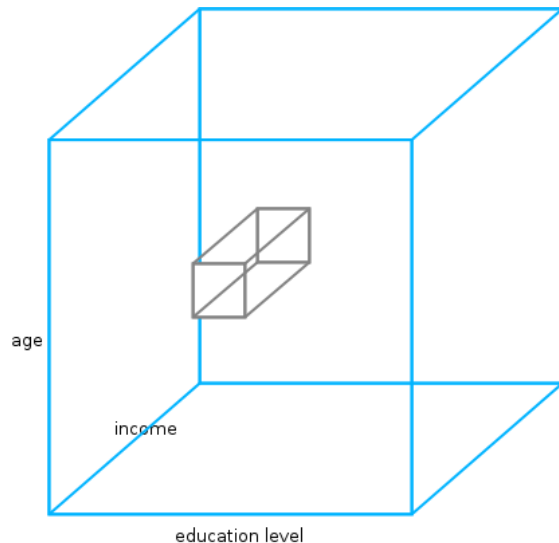
Journey so far



Focused on Technical Innovation



- Multi Dimensional Analysis
 - Not necessarily Geo-Data but any numeric data
 - Query within a Hyper-Cube
 - Map categories to numbers
- e.g.
 - Income, Age, Education level (Bsc = 4, Msc = 5)
 - Timestamp, Log-Level



■ Geo-Data

- GeoJSON: the “Open Standard”
- More complex geometries stored as regions
- Bounding-Box Queries
- e.g. all buildings in San Francisco

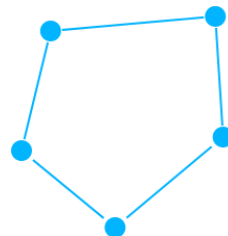


```
{ "type": "Point", "coordinates": [100.0, 0.0] }
```



```
{  
  "type": "LineString",  
  "coordinates": [ [100.0, 0.0], [101.0, 1.0] ]  
}
```

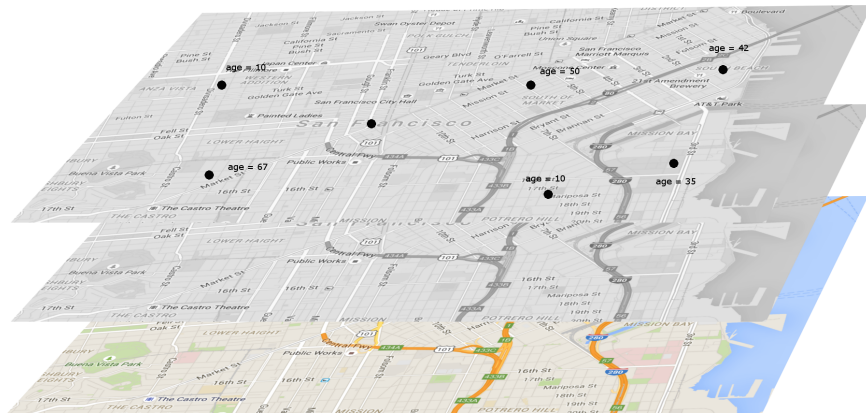
```
{  
  "type": "Polygon",  
  "coordinates": [  
    [ [100.0, 0.0], [101.0, 0.0],  
      [101.0, 1.0], [100.0, 1.0], [100.0, 0.0] ]  
  ]  
}
```



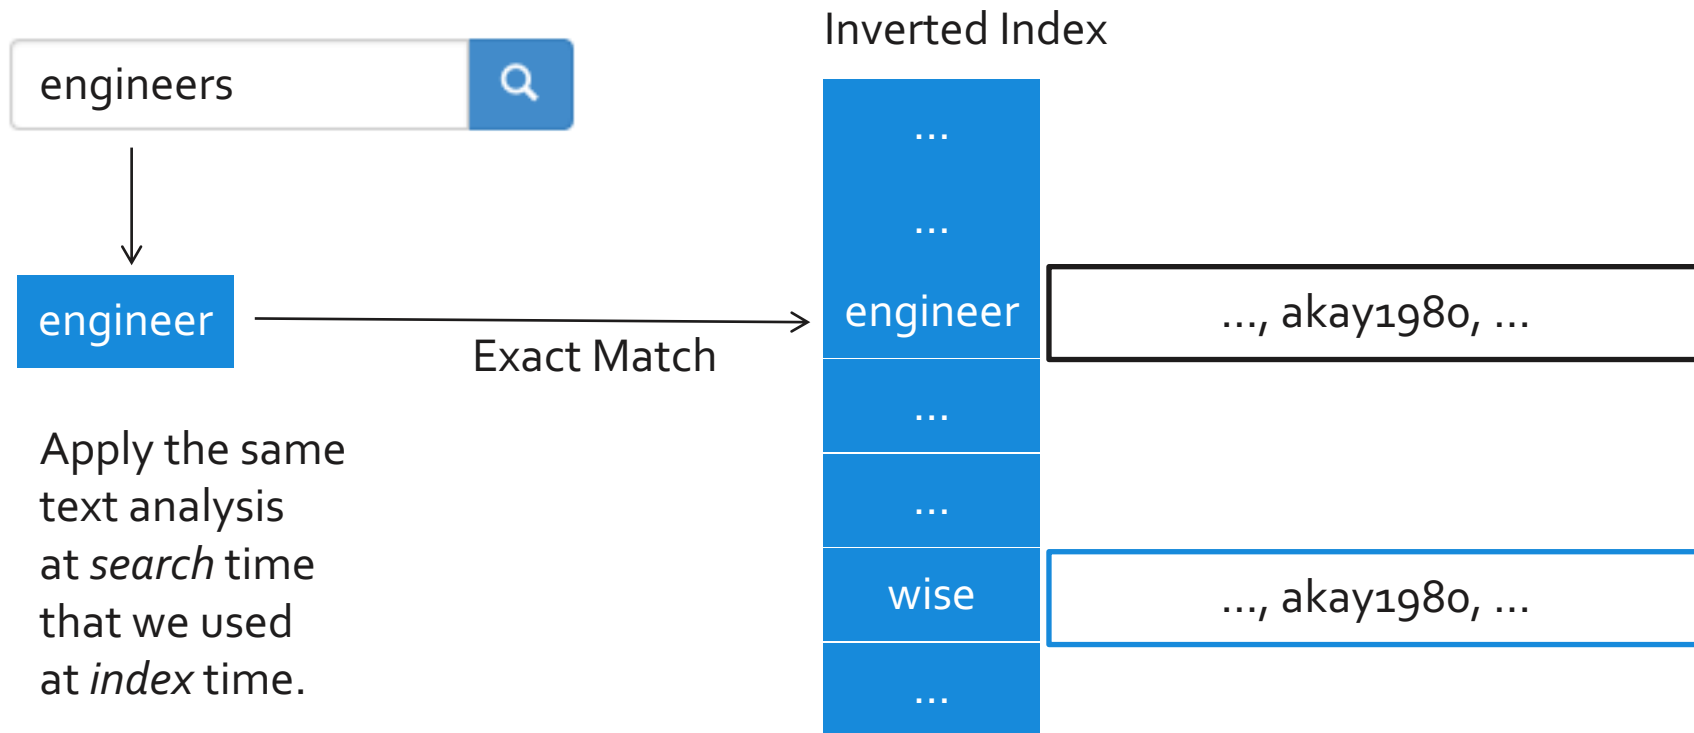
Spatial Indexes



- Combined
 - 2 dimensions for Geo-Data
 - Additional dimensions
 - e.g. all persons with an age greater than 10 in San Francisco



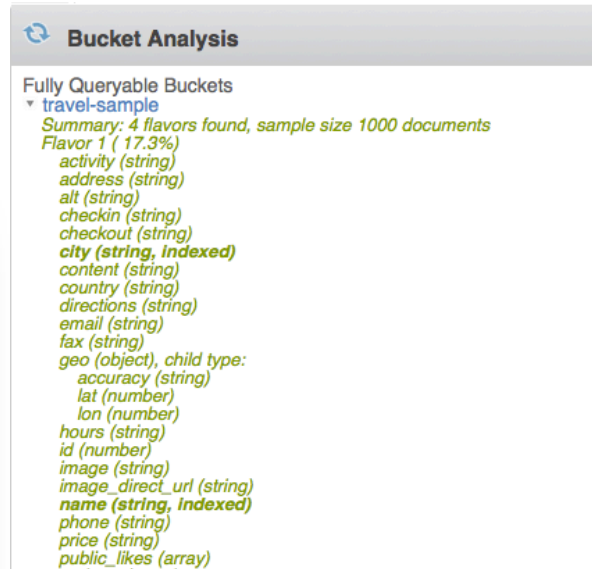
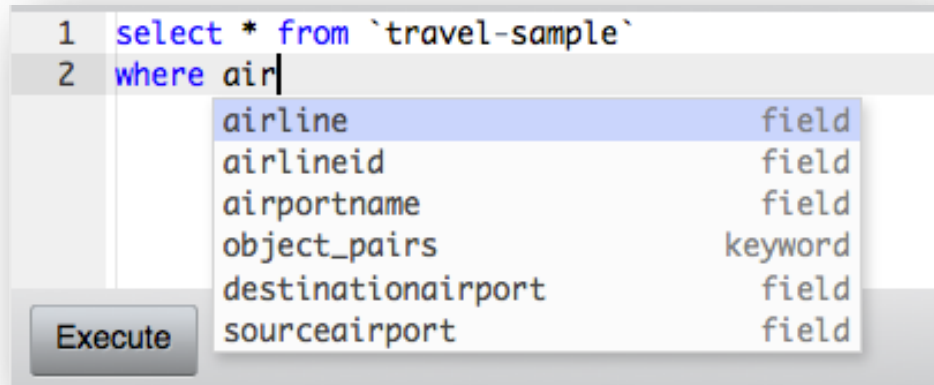
Text Indexes (Developer Preview)



Automatic Schema Inference and Query Editor



- Automatically examines sample of documents from buckets and discovers your schema
 - Document Types and Distribution Stats
 - List of Attributes with Data Types



Memory Optimized Indexes



```
CREATE INDEX idx_Movies ON bucket(  
    DISTINCT ARRAY r.Title FOR r IN Movies  
    END  
);
```

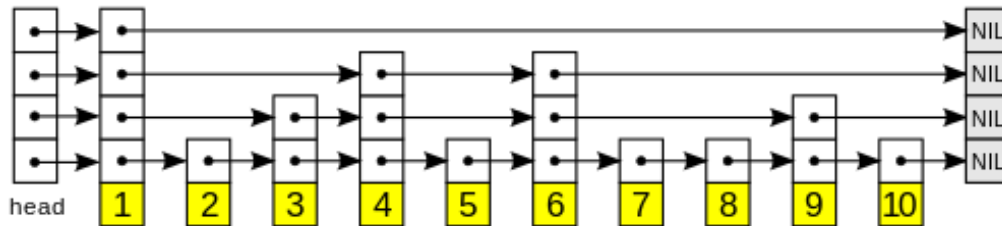
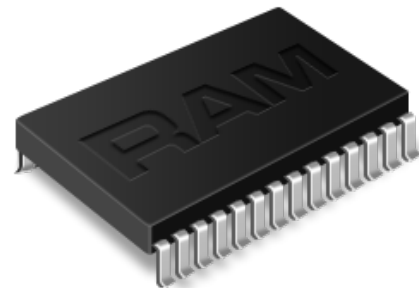
```
SELECT Venue FROM bucket WHERE  
    ANY r IN Movies SATISFIES r.Title="Fight  
Club"  
    END;
```

```
{  
  Start_Date: "1/1/2001",  
  Venue: "AMC 55",  
  Movies:  
  [  
    {Title: "Fight Club",  
      Showtimes:  
      [{Times: ["13:30", "14:45", "21:30"],  
        3D: true},  
       {Times: ["11:30", "15:45", "20:00"],  
        IMAX: true}  
      ...  
      ]  
    },  
    {Title: "Sixth Sense",  
      Showtimes:  
      [{Times: ["10:30", "11:45", "13:30"],  
        3D: true},  
       {Times: ["9:30", "14:45", "20:30"],  
        IMAX: true}  
      ...  
      ]  
    }  
  ]  
}
```

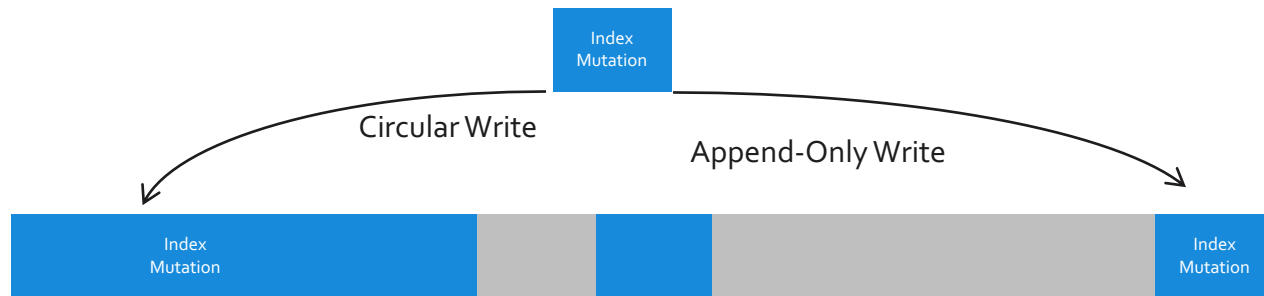
Memory Optimized Indexes



- Optimized for Memory
 - small memory footprint, optimized for lowest latency queries
- Faster Indexing
 - fresh indexes under heavy mutations with lock free index maintenance

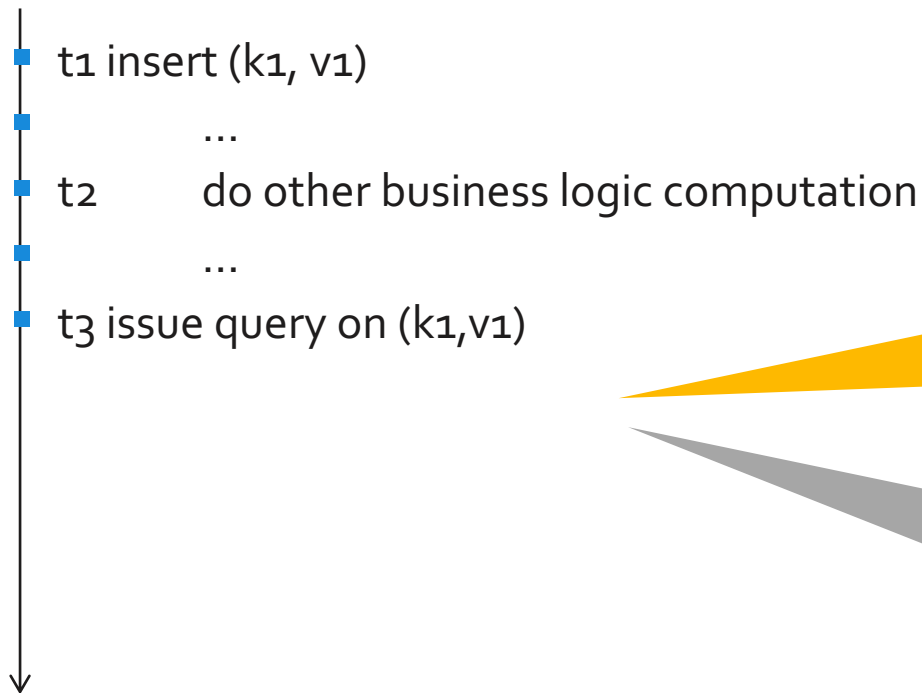


- Reduced Disk IO Requirements
 - Append-Only Writes with frequent full compaction (Version 4.1 & Earlier)
 - Circular-Reuse Writes with reduced full compactions (New in 4.5:)
 - Reused orphaned blocks in the index file
 - Reduce the need for frequent full-compactions of the index file



File for Global Secondary Indexes

High Performance Queries under strict Consistency



RYOW Consistency

Query execution is delayed until all indexes process mutations up to

t1

Strict Request-Time Consistency

(a.k.a stale-false)

Query execution is delayed until all indexes process mutations up to

t3

4.6 Security



- P(lugable) A(uthentication) M(odule)
 - Linux only
- Secret Management
 - Own key chain for Couchbase Server secrets



4.6 Timestamp-based Conflict Resolution for XDCR



- New conflict resolution mode
- Hybrid timestamp as the criteria to order mutations
- Higher hybrid timestamp value as the main factor to determine which document has the most recent mutation
- Hybrid timestamp is replicated across nodes within the same cluster and across clusters
- If both have the same timestamp then revision id is used

4.6 Data Structures



- **MAP:** KV structure, HashMap
- **LIST:** List of objects
- **QUEUE:** Wrapper over a list which offers FIFO semantics
- **SET:** Wrapper over a list which provides the ability to handle unique values

5.0 New UI

